

task-2-meriskill

June 20, 2024

1 Diabetes Prediction Using Machine Learning with Python

1.0.1 Objective:

Develop a machine learning model that can predict whether a patient has diabetes or not based on various medical parameters such as blood glucose levels, insulin levels, BMI, and more.

1.0.2 Key Technologies:

Python: Used for data preprocessing, feature selection, model implementation, and deployment.

Streamlit: Used for creating a user-friendly web interface for diabetes prediction.

Scikit-learn: Used for implementing machine learning algorithms.

Pandas: Used for data manipulation and analysis.

Matplotlib and Seaborn: Used for data visualization.

1.0.3 Benefits:

Early Detection: The model can help detect diabetes at an early stage, enabling timely interventions and improving patient outcomes.

Personalized Medicine: The model can provide personalized predictions based on individual medical parameters, enabling more targeted treatment plans.

Cost Savings: The model can reduce healthcare costs by identifying high-risk patients and enabling preventive measures.

1.0.4 Summary of the dataset:

This dataset seems to contain information related to diabetes risk. The variables include :

* **Pregnancies :** The number of pregnancies an individual has had.

* Glucose : Glucose levels in the blood.

* BloodPressure : Blood pressure readings.

* SkinThickness : Thickness of a skinfold at a certain location on the body.

* Insulin : Levels of insulin in the blood.

* BMI (BODY Mass Index) : A measure of body fat based on height and weight.

* DiabetesPredictionFunction: A Function that scores the likelihood of diabetes based on family history.

* Age: Age of the individuals

* Outcome: abinary variable indicating the presence (1) or absence(0) of a diabetes outcome.

```
[2]: #Importing the required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[3]: #Load the Dataset
dia=pd.read_csv("D:\chrome (downloads)\Project 2 MeriSKILL\diabetes.csv")
```

```
[4]: dia
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1

3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[18]: pip install matplotlib
```

```
Requirement already satisfied: matplotlib in c:\users\venky\anaconda3\lib\site-
packages (3.8.0)
Requirement already satisfied: contourpy>=1.0.1 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (1.2.0)
Requirement already satisfied: cyclor>=0.10 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (1.4.4)
Requirement already satisfied: numpy<2,>=1.21 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (1.26.4)
Requirement already satisfied: packaging>=20.0 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (23.1)
Requirement already satisfied: pillow>=6.2.0 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (10.2.0)
Requirement already satisfied: pyparsing>=2.3.1 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in
c:\users\venky\anaconda3\lib\site-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\venky\anaconda3\lib\site-
packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

dataset outline: This contains 768 observations with 8 feature columns and a target variable 'outcome'

```
[4]: #checking the outline of the dataset
dia.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
```

```

0    Pregnancies      768 non-null    int64
1    Glucose          768 non-null    int64
2    BloodPressure    768 non-null    int64
3    SkinThickness    768 non-null    int64
4    Insulin          768 non-null    int64
5    BMI              768 non-null    float64
6    DiabetesPedigreeFunction  768 non-null    float64
7    Age              768 non-null    int64
8    Outcome          768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[5]: #checking for null values
dia.isnull().sum()
```

```

[5]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64

```

```
[6]: #Checking the dupilcate
dia.duplicated().sum
```

```

[6]: <bound method Series.sum of 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      763     False
      764     False
      765     False
      766     False
      767     False
      Length: 768, dtype: bool>

```

```
[7]: dia.duplicated().sum()
```

```
[7]: 0
```

```
[8]: #Analysing the summary of the dataset
dia.describe()
```

```
[8]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000


```
[9]: dia['Outcome'].value_counts()
```

```
[9]: Outcome
0      500
1      268
Name: count, dtype: int64
```

1.1 Exploratory Data Analysis

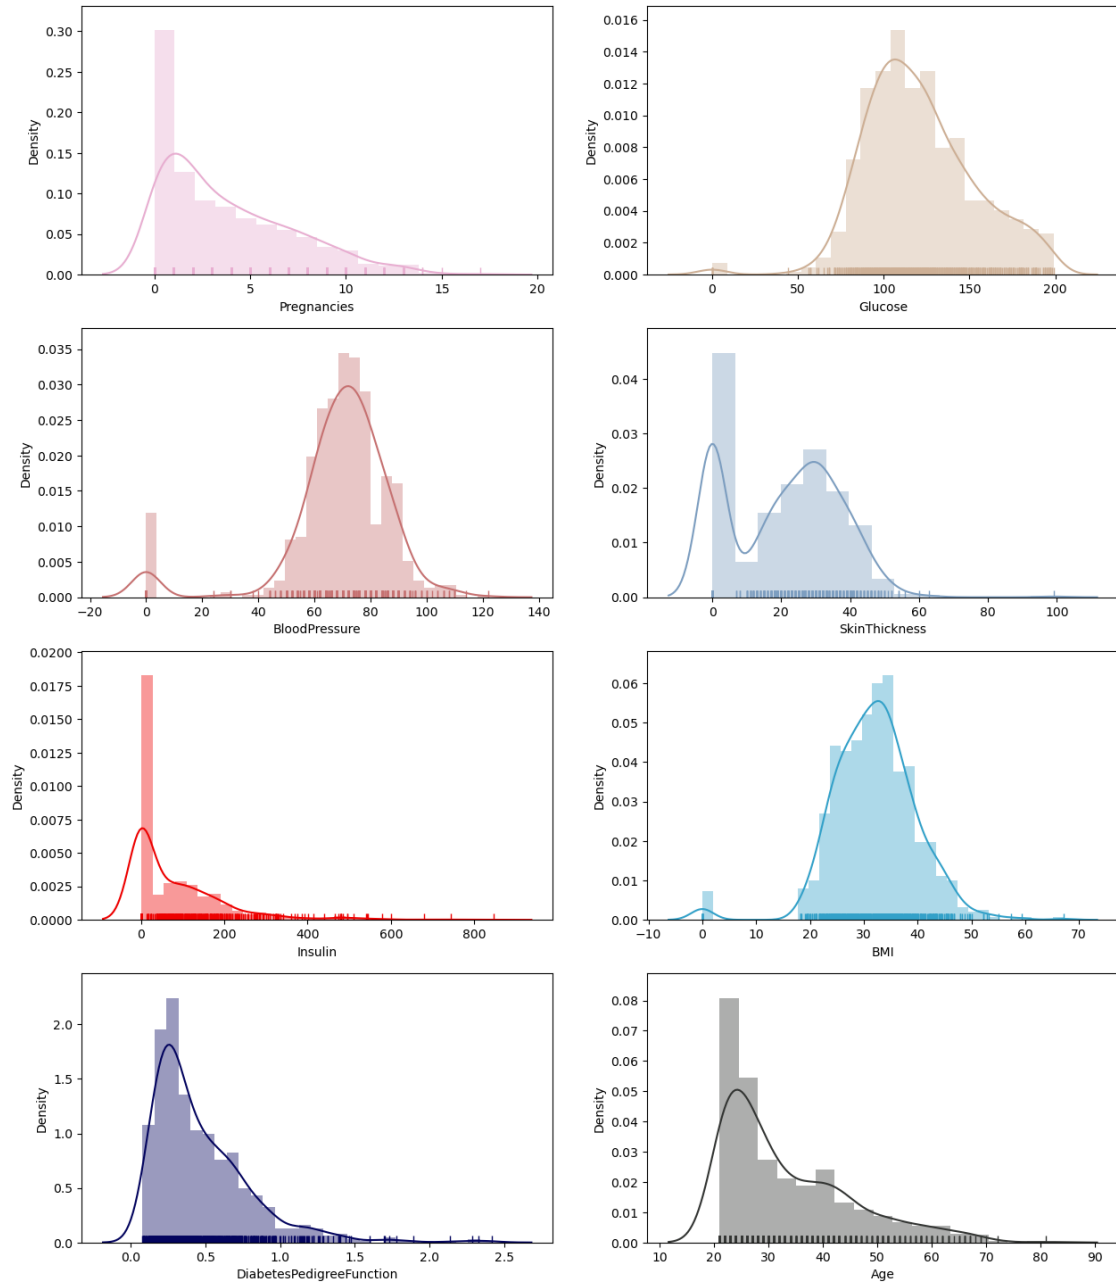
```
[10]: #visualizing bar graph of the outcome
# 1 means diabetes patient and 0 means no diabetes patient
dia['Outcome'].value_counts().plot(kind='bar', color= ['red','blue'])
```

```
[10]: <Axes: xlabel='Outcome'>
```



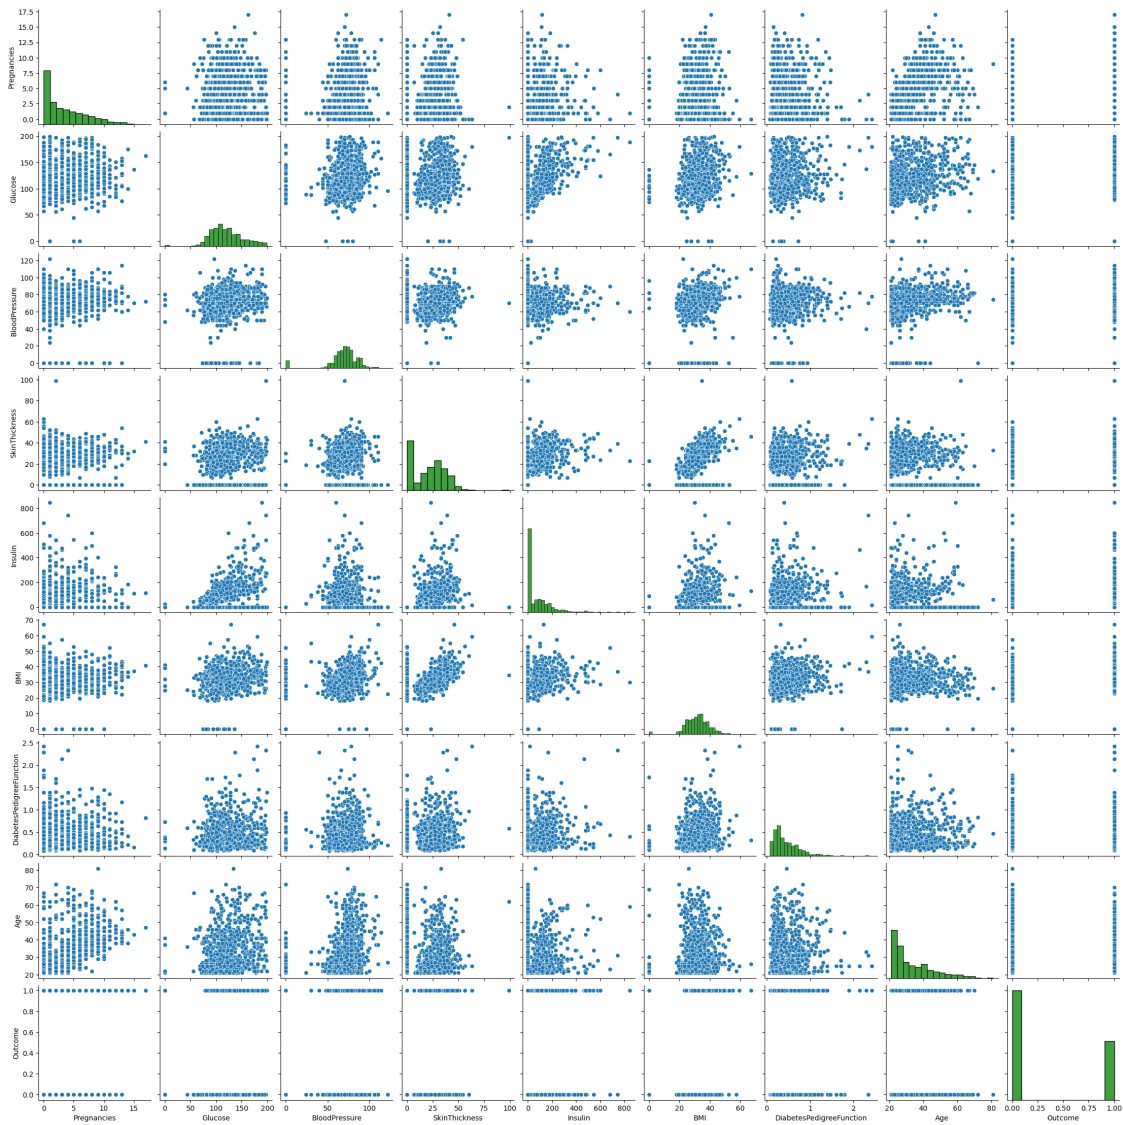
```
[25]: fig, axs = plt.subplots(4, 2, figsize=(15, 18))
      axs = axs.flatten()
      sns.distplot(dia['Pregnancies'], rug=True, color='#E8AED1', ax=axs[0])
      sns.distplot(dia['Glucose'], rug=True, color='#CDAF95', ax=axs[1])
      sns.distplot(dia['BloodPressure'], rug=True, color='#C67171', ax=axs[2])
      sns.distplot(dia['SkinThickness'], rug=True, color='#7D9EC0', ax=axs[3])
      sns.distplot(dia['Insulin'], rug=True, color='#EE0000', ax=axs[4])
      sns.distplot(dia['BMI'], rug=True, color='#33A1C9', ax=axs[5])
      sns.distplot(dia['DiabetesPedigreeFunction'], rug=True, color='#03045e',
      ↪ax=axs[6])
      sns.distplot(dia['Age'], rug=True, color='#333533', ax=axs[7])
```

```
[25]: <Axes: xlabel='Age', ylabel='Density'>
```




```
[28]: sns.pairplot(dia , diag_kws={'color':'green'})
```

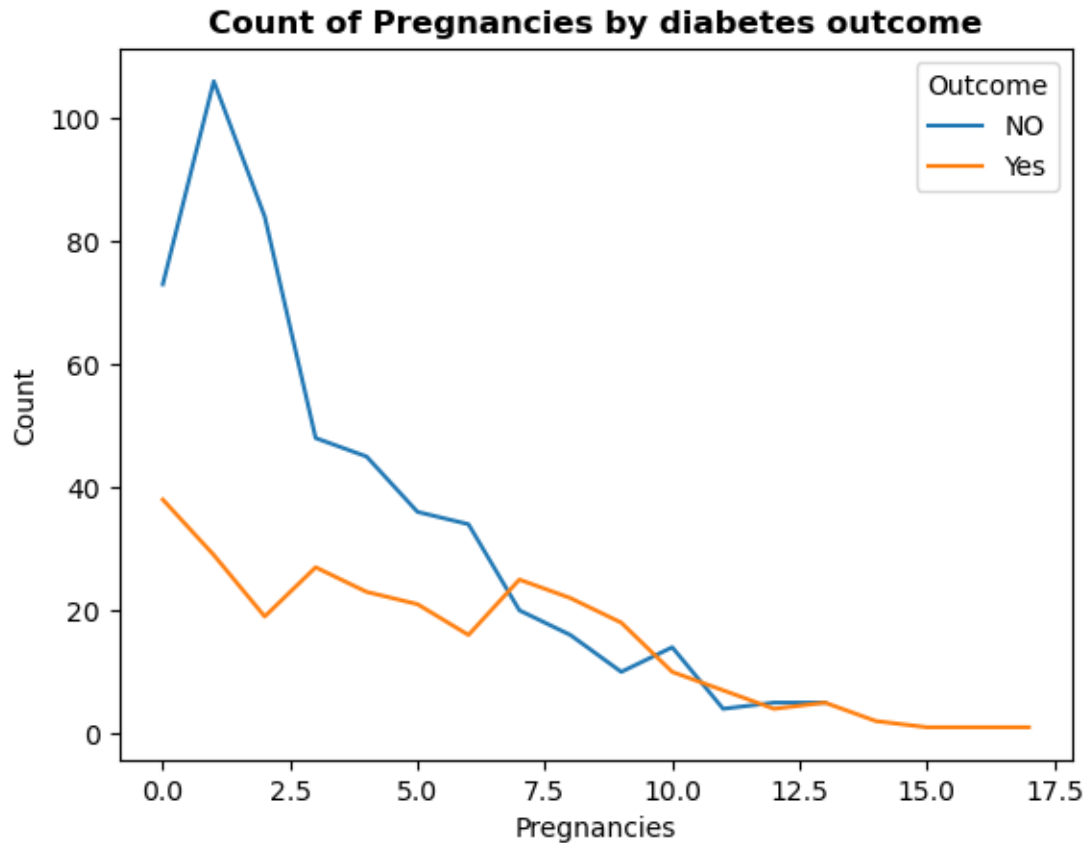
```
[28]: <seaborn.axisgrid.PairGrid at 0x26827d2dc90>
```



```
[31]: # Visualizing HeatMap
sns.heatmap(dia.corr(), annot=True ,fmt='.2f')
plt.title('Features Vs Features')
plt.xticks(rotation=45)
plt.show()
```

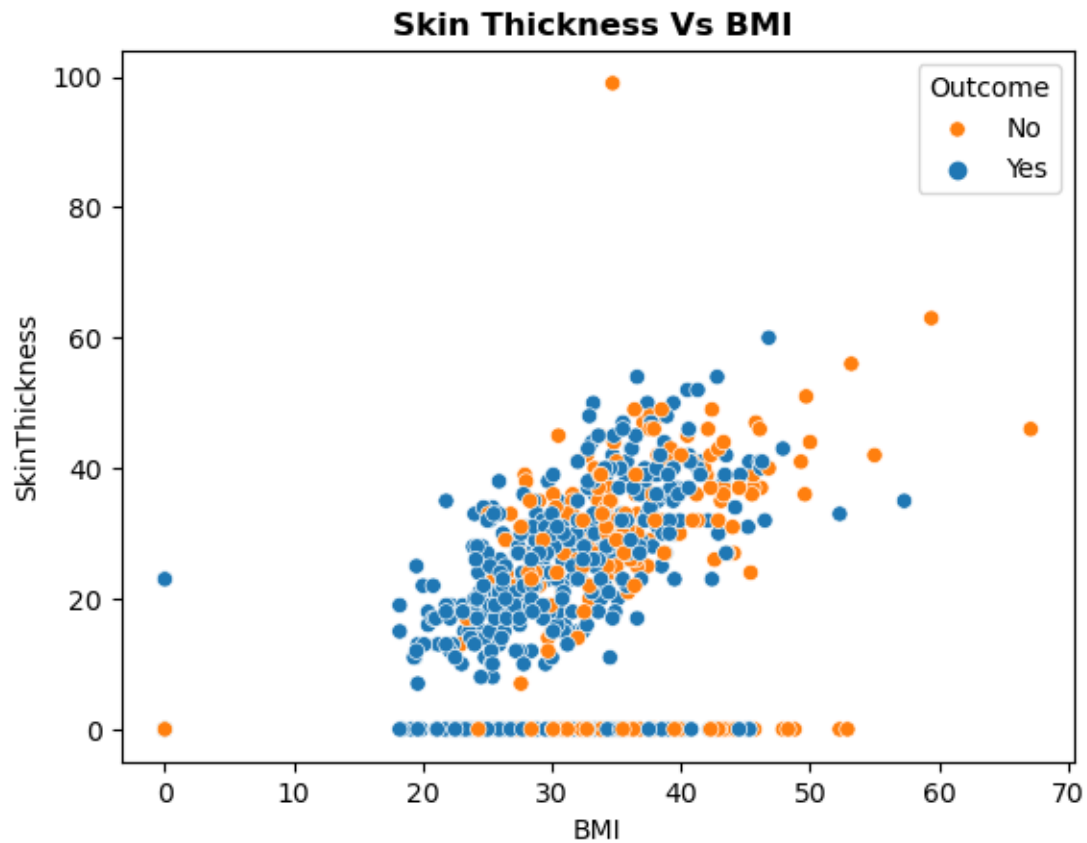



```
[40]: dia.groupby(['Pregnancies', 'Outcome']).size().unstack(level=1).plot(kind='line')
plt.ylabel('Count')
plt.legend(title='Outcome', labels= ['NO', 'Yes'])
plt.title('Count of Pregnancies by diabetes outcome', weight='bold')
plt.show()
```




```
[41]: #Insight:
      # The number of having diabetesbis less when the number of pregnancies is low
      # The Possibility of having diabetes increas as the number of pregnancies is low
```


```
[43]: sns.scatterplot(data=dia , x='BMI' , y='SkinThickness' , hue='Outcome')
      plt.legend(title= 'Outcome', labels=['No', 'Yes'])
      plt.title('Skin Thickness Vs BMI ', weight ='bold')
      plt.show()
```




```
[48]: sns.scatterplot(data=dia , x='Glucose' , y='BloodPressure' , hue='Outcome')
plt.legend(title= 'Outcome', labels=['No', 'Yes'])
plt.title('Blood Pressue Vs Glucose ', weight ='bold')
plt.show()
```




```
[49]: sns.scatterplot(data=dia , x='Age' , y='BloodPressure' , hue='Outcome')
plt.legend(title= 'Outcome', labels=['No', 'Yes'])
plt.title('Age Vs BloodPressure ', weight ='bold')
plt.show()
```



```
[50]: sns.scatterplot(data=dia , x='Age' , y='BMI' , hue='Outcome')  
plt.legend(title= 'Outcome', labels=['No', 'Yes'])  
plt.title('Age Vs BMI ', weight = 'bold')  
plt.show()
```



```
[51]: sns.scatterplot(data=dia , x='Age' , y='DiabetesPedigreeFunction' ,
    ↪ hue='Outcome')
plt.legend(title= 'Outcome', labels=['No', 'Yes'])
plt.title('Age Vs DiabetesPedigreeFunction ', weight ='bold')
plt.show()
```



```
[5]: from sklearn import metrics
      from sklearn.model_selection import train_test_split
```

```
[6]: #splitting data into features and target
      x= dia.drop(["Outcome"], axis= "columns")
      y= dia["Outcome"]
```

```
x.head()
```

```
[7]: y.head()
```

```
[7]: 0    1
      1    0
      2    1
      3    0
      4    1
      Name: Outcome, dtype: int64
```

```
[8]: #splitting the dataset into training and test set
```

```
x_test, x_train, y_test, y_train = train_test_split(x,y, test_size= 0.2,
↳train_size= 0.8,random_state=123,shuffle= True)
```

```
[9]: x_test.shape,x_train.shape, y_test.shape, y_train.shape
```

```
[9]: ((614, 8), (154, 8), (614,), (154,))
```

2 Model building

2.0.1 Logistic Regression

```
[10]: #import Logistic regression from sklearn module
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
↳accuracy_score
```

```
[19]: LR=LogisticRegression()
```

```
[21]: logistic_model = LogisticRegression()
logistic_model.fit(x_train,y_train)
```

```
[21]: LogisticRegression()
```

```
[25]: LR.fit(x_train, y_train)
```

```
[25]: LogisticRegression()
```

```
[26]: y_pred= LR.predict(x_test)
```

```
[27]: accuracy= accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.754071661237785

```
[28]: CM = confusion_matrix(y_test, y_pred)
print(CM)
```

```
[[344  60]
 [ 91 119]]
```

```
[29]: class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.79	0.85	0.82	404
1	0.66	0.57	0.61	210

accuracy			0.75	614
macro avg	0.73	0.71	0.72	614
weighted avg	0.75	0.75	0.75	614

3 Decision Tree

```
[32]: from sklearn.tree import DecisionTreeClassifier
```

```
[33]: DT= DecisionTreeClassifier()
```

```
[34]: DT.fit(x_train, y_train)
```

```
[34]: DecisionTreeClassifier()
```

```
[36]: y_pred = DT.predict(x_test)
```

```
[37]: DT.score(x_test, y_test)
```

```
[37]: 0.6775244299674267
```

```
[38]: accuracy_DT = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_DT)
```

```
Accuracy: 0.754071661237785
```

```
[39]: CM_DT = confusion_matrix(y_test, y_pred)
print(CM_DT)
```

```
[[344  60]
 [ 91 119]]
```

```
[40]: class_report= classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.79	0.85	0.82	404
1	0.66	0.57	0.61	210
accuracy			0.75	614
macro avg	0.73	0.71	0.72	614
weighted avg	0.75	0.75	0.75	614

4 Random Forest Classifier

```
[41]: from sklearn.ensemble import RandomForestClassifier
```

```
[42]: random_forest = RandomForestClassifier(n_estimators=10)
```

```
[43]: random_forest.fit(x_train, y_train)
```

```
[43]: RandomForestClassifier(n_estimators=10)
```

```
[44]: y_pred = random_forest.predict(x_test)
```

```
[45]: random_forest.score(x_test, y_test)
```

```
[45]: 0.744299674267101
```

```
[46]: accuracy_RF = accuracy_score(y_test, y_pred)
      print("Accuracy:", accuracy_RF)
```

Accuracy: 0.744299674267101

```
[48]: CM_RF = confusion_matrix(y_test, y_pred)
      print(CM_RF)
```

```
[[344  60]
 [ 97 113]]
```

```
[49]: class_report = classification_report(y_test, y_pred)
      print(class_report)
```

	precision	recall	f1-score	support
0	0.78	0.85	0.81	404
1	0.65	0.54	0.59	210
accuracy			0.74	614
macro avg	0.72	0.69	0.70	614
weighted avg	0.74	0.74	0.74	614

5 Gaussian Naive Bayes Classifier

```
[14]: from sklearn.naive_bayes import GaussianNB
      naive_bayes = GaussianNB()
```

```
[15]: naive_bayes.fit(x_train, y_train)
```

```
[15]: GaussianNB()
```

```
[16]: y_pred_NB = naive_bayes.predict(x_test)
```

```
[54]: naive_bayes.score(x_test, y_test)
```

```
[54]: 0.741042345276873
```

```
[17]: accuracy_NB = accuracy_score(y_test, y_pred_NB)
      print("Accuracy:", accuracy_NB)
```

```
Accuracy: 0.741042345276873
```

```
[18]: CM_NB = confusion_matrix(y_test, y_pred_NB)
      print(CM_NB)
```

```
[[329  75]
 [ 84 126]]
```

```
[19]: class_report = classification_report(y_test, y_pred_NB)
      print(class_report)
```

	precision	recall	f1-score	support
0	0.80	0.81	0.81	404
1	0.63	0.60	0.61	210
accuracy			0.74	614
macro avg	0.71	0.71	0.71	614
weighted avg	0.74	0.74	0.74	614

6 Interpretation and summary report

- Imported the necessary modules for the project and checked the data
- Performed Exploratory analysis to visualize the distribution of different features
- Performed preprocessing steps to treat the missing values
- Used 4 Models (Logistic Regression , random Forest , Decision tree and Naive Bayes Classifier) to find the best model for prediction
- Evaluated the performance using the Accuracy, Precision, Recall, and F1 score
- Based on the performance evaluation, Logistic Regression performed well in predicting if someone has diabetes or not
- Logistic regression has the highest accuracy of 76%
- This model also exhibited reasonable precision and recall, indicating its ability to correctly classify both positive and negative cases of diabetes.
- So, The Machine learning approach , specifically the Logistic regression , can be a valuable tool for predicting diabetes outcomes based on health-related variables

[]: